

PROGRESS REPORT ON STUDY OF NYC TRAFFIC

VAIBHAV KARVE

1. RESTRICTING TO A BLOCK

We restrict our attention to the block between W. 44th to 45th St. and 8th to 9th Ave. This block was selected at random. On extracting the data from the data file for 2011, we see that these streets are all one-way access only.

1.1. Naming convention.

1.1.1. *Intersections.*

- A = W. 44th St. and 8th Ave. (Node_id = 42435671)
- B = W. 44th St. and 9th Ave. (Node_id = 42443561)
- C = W. 45th St. and 8th Ave. (Node_id = 42432700)
- D = W. 45th St. and 9th Ave. (Node_id = 42432703)

1.1.2. *Roads.*

- BA (Link_id = 128255)
- AC (Link_id = 169017)
- CD (Link_id = 182993)
- DB (Link_id = 181188)

We restrict our attention to only these links, one at a time. For each link, from the database, we extract 2 separate arrays: one giving the average travel time in seconds, for every hour of the year; and the other giving the number of trips on that used that link, for every hour of the year.

1.2. Periodicity analysis.

1.2.1. *Stratifying the data.* Intuitively, one may expect that traffic patterns repeat themselves every 7 days (weekly) or maybe every 30 days (monthly). Or perhaps, there is no such periodicity. Whatever be the case, the periodicity should not be imposed, but rather should be inferred from the data itself. To do so, we assume a particular period in days, call it **period** and divide the entire data into $24 \times \text{period}$ number of bins. This converts our data from flat lists of trips and traveltimes to something that may be viewed as stratified data.

Stratified data now looks like:

$$\begin{pmatrix} \text{Stratum}_1 & \text{Stratum}_2 & \text{Stratum}_3 & \cdots & \text{Stratum}_N \\ 5 & 2 & 32 & \cdots & 10 \\ ? & 7 & 5 & \cdots & ? \\ 33 & ? & 12 & \cdots & ? \\ \vdots & \vdots & \vdots & \cdots & \vdots \end{pmatrix}$$

Here, $N = 24 \times \text{period}$. We let period range from 2 to 38 for our periodicity analysis.

Date: June 3, 2016.

1.2.2. *Missing values.* The question marks (?) in the above matrix correspond to those hours for which we have no data on our link. This ofcourse does not mean that there is no traffic on that link at that time, just that we don't know what it is. These are to be treated as missing values in our data and need to be somehow inferred.

The links BA , AC , CD and DB have 20%, 0%, 0% and 1% of their values missing, respectively.

The most natural first approximation for these missing values is the mean value of the corresponding stratum to which each missing value belongs.

The Inferred Stratified data looks like:

$$\begin{pmatrix} \text{Stratum}_1 & \text{Stratum}_2 & \text{Stratum}_3 & \cdots & \text{Stratum}_N \\ 5 & 2 & 32 & \cdots & 10 \\ \text{Mean}_1 & 7 & 5 & \cdots & \text{Mean}_N \\ 33 & \text{Mean}_2 & 12 & \cdots & \text{Mean}_N \\ \vdots & \vdots & \vdots & \cdots & \vdots \end{pmatrix}$$

Here, we calculated the mean for each stratum by ignoring the missing values.

1.2.3. *Stratified variances.* To establish the optimal choice of **period** for the data, we calculate the variance of the inferred stratified data we obtained in the previous subsection.

$$\text{Variance of stratified data} \approx \frac{1}{n} \sum_n \text{Variance of each stratum}$$

where $n = 24 \times \text{period}$ is the total number of strata.

If there truly does exist a periodicity in the data, for the correct **period**, the values within each stratum will lie close to each other and hence, the variance will be minimized. Hence, we look for dips in the variance.

1.2.4. *Conclusion of periodicity analysis.* We calculate the variance for each value of **period** from 2 to 38. The graph of the variances for link BA for the No. of Trips during 2011 is given below:




FIGURE 1. The dips in variance correspond to **period** = 7, 14, 21, 35.

The other three links show similar dips in variance at **period** values which are multiples of 7, in both data – Travel times as well as No. of Trips. **Conclusion:** NYC traffic data has a 7-day periodicity.

1.3. Periodicity analysis via autocorrelation. An easier, and more standard way to establish periodicity would be to calculate the autocorrelation of our data for four links assuming different lags (**period** values) of 1 to 39 days.

For a discrete set of observations $\{X_1, X_2, \dots, X_n\}$, an estimate of the autocorrelation can be obtained as,

$$\widehat{\text{Autocorrelation}}(k) = \frac{1}{(n-k)\sigma^2} \sum_{t=1}^{n-k} (X_t - \mu)(X_{t+k} - \mu)$$

where $k = 24 \times \text{period}$ is the lag in hours for which we are testing autocorrelation in our data, and μ and σ^2 are the mean and variance of the data respectively.

1.3.1. Missing values. We calculate the mean (μ) of the data, ignoring the missing values. We then replace the missing values with μ . We calculate the variance (σ^2) of this *corrected* data. We then proceed to calculate the autocorrelation for different values of **period** using the formula mentioned in the previous subsection.

1.3.2. Conclusion of autocorrelation study. We plot the autocorrelations for Trips data and Traveltimes separately, for all four road links simultaneously. The autocorrelation is high for **period** values that are a multiple of 7. This confirms what was seen by the statifying the data – NYC traffic patterns repeat every seven days.




FIGURE 2. Autocorrelations for Travel-times data of all four road links. The peaks in autocorrelation correspond to **period** = 7, 14, 21, 35.



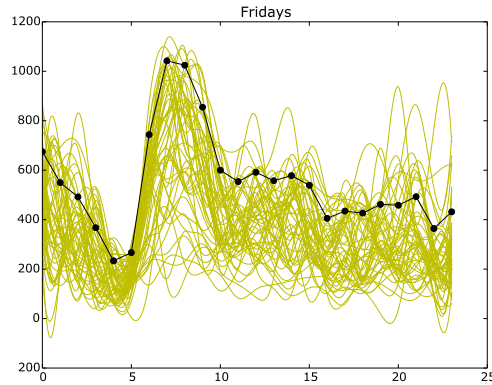
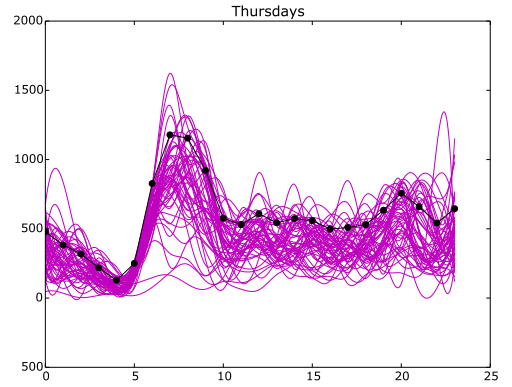





FIGURE 3. Autocorrelations for Trips data of all four road links. The peaks in autocorrelation correspond to $\text{period} = 7, 14, 21, 35$.

1.4. Visualizing the data. Keeping in view the conclusions drawn from the autocorrelation analysis, henceforth, we try to analyze each dat of the week seperately. We try to visualize the data by plotting different parts of the data file, in the hopes that it will point towards directions in which we can carry out further study.

1.4.1. Daily pattern. Below are shown the graphs for each day of the week. These plots are also overlaid with the $r = 1$ results from Non-negative Matrix factorization. NMF is discussed in greater detail in the next section.





Daily trends in Trips data on link AC for each day of the week, for each week in 2011. The trends are overlaid with the factor W as calculated using Non-negative Matrix Factorization of the data, with $r = 1$.

As can be seen from the plots, the rank-one factor of our data matrix can act as a low-rank approximation.

Unresolved issue: Why is the W vector slightly higher than the “mean” daily trend?

1.4.2. *Observations from visual examination of daily trends.* From visual examination of the graphs showing daily trends of traffic, we can make a few observations:

- (1) After superimposing the plots for each day of the week, one can clearly see a daily pattern. For example, on an average Monday:
 - Midnight to 1am: the traffic drops to a minimum value at 1am, settling down for the night.
 - 1am to 5am: very thin traffic.
 - 5am to 10am: very high traffic. Traffic peaks at 7am.

- 10am to midnight: traffic plateaus at a value less than rush hour. This region has very high variance across all Mondays i.e. the exact traffic in this time-period varies a lot from Monday-to-Monday.
- (2) The schedule outlined above holds for all the weekdays, i.e. Monday-Friday behave the same.
 - (3) The patterns for both Saturday and Sunday are very similar to each other, but different from the weekday schedules:
 - Midnight to 3am: high traffic, with high variance across all weekends. Traffic peaks at 2am.
 - 3am to 5am: traffic drops drastically and settles at a minimum value at 5am.
 - 5am to 9am: very thin traffic.
 - 9am to midnight: traffic plateaus at a value less than rush hour. This region has very high variance across all weekends i.e. the exact traffic in this time-period varies a lot from weekend-to-weekend.
 - (4) On calculating the derivative matrix for each day of the week, one can see that those plots show similar weekday vs. weekend patterns. (Graphically though, the derivative matrix's plot is more difficult to interpret compared to the actual transport matrix's plot.)
 - (5) Regions with high variance: As pointed out in 1. and 3., certain periods have a high variance even after taking into consideration the weekly periodicity in data. These might point to other factors that need to be considered eg: seasonal variations in traffic.
 - (6) All the above analysis was done with data from a single road link. However, other links show similar (but not completely identical) patterns.
 - (7) Item 6. points to what might be the most important and counter-intuitive aspect of all: these patterns in traffic are present not just globally but even locally. One might expect that even though we may see some small patterns when studying a single link, more complex phenomena like rush-hour, weekend vs. weekday patterns etc. would not be visible at the street-level and would show up only once we start considering larger networks of roads. However, items 1. to 4. go towards showing that these phenomena percolate down to the street level!

1.5. Non-negative Matrix Factorization.

1.5.1. *The theory.* Suppose we have a large matrix D of datapoints which we wish to understand. It is easier to handle this data, and make sense of it if the dimension can somehow be reduced. One of several matrix factorization methods can be used to achieve this dimensionality reduction.

Given a $n \times m$ -matrix D , we wish to find matrices W and H of sizes $n \times r$ and $r \times m$ respectively satisfying

$$D \approx WH$$

Here, the choice of r is up to us. Matrix factorization is useful only when we choose $r < \min\{n, m\}$. There exist several algorithms, starting from an initial guess of W and H , iteratively update these matrices, to guarantee convergence to D in a finite number of steps.

Moreover, since all our datapoints (i.e. matrix entries) either consist of road-links, average speeds, or average travel times, in order to make sense of the matrix factors, it is desirable to constrain the elements of W and H to always be non-negative. This can be guaranteed by the choice of iterative updation rules. We choose the algorithm called Non-negative Matrix Factorization (NMF) for this very reason.

1.5.2. The algorithm.

- (1) Choose a value for $r < \min\{n, m\}$, where $n \times m$ is the dimension of D .
- (2) Initialize W and H as random matrices, with sizes $n \times r$ and $r \times m$ respectively.
- (3) Define error as

$$\begin{aligned} \text{Error} &= \|D - WH\|^2 \\ &= \sum_{i=1}^n \sum_{j=1}^m (D_{ij} - (WH)_{ij})^2 \end{aligned}$$

- (4) In order to find a local minima of Error, iteratively update W and H using the following rules:

$$W_{ij} \leftarrow W_{ij} \frac{(DH^T)_{ij}}{(WHH^T)_{ij}}$$

and

$$H_{ij} \leftarrow H_{ij} \frac{(W^T D)_{ij}}{(W^T W H)_{ij}}$$

- (5) Stop when Error gets smaller than a chosen tolerance level.

1.5.3. Implementing the algorithm. For the sake of explaining how we implement NMF in our dataset, we restrict our attention to only the Trips data for Link AC in the year 2011. This can be written out as a single vector of size 8760, starting at the trips data for January 1, 12am-1am and ending December 31, 11pm -12am.

In view of the autocorrelation results showing us a weekly-periodicity in traffic data, we can stack the Trips data such that each week is respresented in a separate column. The resulting matrix, which we call D looks like this:

$$D = \begin{matrix} & \text{Week 1} & \text{Week 2} & \cdots & \text{Week 53} \\ \begin{matrix} \text{Sat 0000-0100 hrs} \\ \text{Sat 0100-0200 hrs} \\ \vdots \\ \text{Fri 2200-0300 hrs} \\ \text{Fri 2300-0000 hrs} \end{matrix} & \begin{pmatrix} * & * & \cdots & * \\ * & * & \cdots & * \\ \vdots & \vdots & & \vdots \\ * & * & \cdots & 0 \\ * & * & \cdots & 0 \end{pmatrix} \end{matrix}$$

Note that the trailing values in the last column will need to be set to zero because a year does not contain 53 full weeks and there will be data points “missing” in the last week.

1.5.4. Conclusions from NMF. Implementing the NMF algorithm for $r = 1$ yeilds a vector W which accounts for the data to a great extent, as seen in the in the graphs showing trends.

1.6. Unresolved issues and things that need further investigation.

- (1) Regions with high variance: As pointed out in 1. and 3., certain periods have a high variance even after taking into consideration the weekly periodicity in data. These might point to other factors that need to be considered eg: seasonal variations in traffic.
- (2) In the graphs, the W vector is slightly higher than the “mean” daily trend, i.e. W seems to be enveloping the traffic trend rather than approximating it. The reason for this is unclear.
- (3) While we have made use of the $r = 1$ NMF, it is not clear how the results from higher-rank factorizations is to be interpreted.

- (4) The NMF analysis outlined above was done only for link AC because AC has almost no missing data. We need to figure out a way to fill in the missing data before applying NMF to other links. Or perhaps, we need to modify the NMF algorithm to infer the missing datapoints from the factors themselves. (Weighted matrix entries need to be explored as a way of doing this.)